

# Intro to ML-ops

## From Your PC to the Cloud

Sveinung Myhre

ReLU NTNU

September 30, 2025

# Outline

- 1 Introduction
- 2 Docker Containers
- 3 Kubernetes
- 4 GPU Glossary
- 5 Hands-on Coding
- 6 Questions & Discussion

Welcome to this brief introduction to ML-ops!

**Today's mission:**

*“Acquire the knowledge to put a model from your PC into the cloud”*

## Assumptions:

- You know Git basics
- You're familiar with basic Linux commands
- You know what a REST API looks like

**If not, that's OK!** These are concepts you'll become familiar with throughout your tech career.

## Mixed Experience Levels

For some: this will be elementary

For others: entirely new

**Everyone:** please come with questions!

**Goal:** Acquire the knowledge to put a model from your PC into the cloud

**Topics we'll cover:**

- Typical ML-ops pipeline overview
- Docker & Containerization
- Kubernetes & Orchestration
- Cloud Deployment strategies
- Hands-on coding session

# Typical ML-ops Pipeline

**ML-ops workflow:** Development → Testing → Deployment → Monitoring

Code → Container → Orchestration → Cloud → Production

**Today we focus on:** CI/CD components

- Docker containers
- Kubernetes orchestration
- Auto-scaling clusters
- Reverse proxy architecture

# Why I'm Teaching This: Docker

## What is Docker?

- Containerization platform
- Package applications with dependencies
- Consistent environments across machines

## Why is it important?

- Eliminates "works on my machine" problems
- Essential for modern cloud deployments
- Industry standard for ML model deployment

# Evolution: Machine $\rightarrow$ VM $\rightarrow$ Container

## Physical Machine

- Heavy resource usage
- OS dependency
- Difficult scaling

## Virtual Machine

- Full OS overhead
- Better isolation
- Still resource heavy

## Container

- Shared kernel
- Lightweight
- Fast startup

## Contrast: Docker containers vs Python virtual environments

- venv: Python packages only
- Docker: Entire runtime environment



# Essential Docker Commands

## Most important commands to know:

- `docker build -t myapp .` - Build image
- `docker run myapp` - Run container
- `docker ps` - List running containers
- `docker images` - List available images
- `docker pull ubuntu` - Download image
- `docker stop <container-id>` - Stop container

## Key files:

- `Dockerfile` - Build instructions
- `.dockerignore` - Files to exclude

**Container registries** store and distribute Docker images

**Popular services:**

- **Docker Hub** - Free public registry
- **AWS ECR** - Amazon's container registry
- **Google Container Registry**
- **GitHub Container Registry**

**Analogy:** Docker Hub is to Docker as GitHub is to Git

- Git/Docker: Version control tools
- GitHub/Docker Hub: Cloud hosting services

## Time Constraint Strategy

Instead of traditional pair-programming:

**Our approach:**

- ① **Run the code** - Execute provided examples
- ② **Read what it does** - Understand the implementation
- ③ **Explain to each other** - Discuss with peers
- ④ **Ask AI assistants** - Use ChatGPT/LLMs for clarification

**Goal:** Maximum learning in minimal time!

# Why I'm Teaching This: Kubernetes

## What is Kubernetes?

- Container orchestration platform
- Manages containerized applications at scale
- Handles deployment, scaling, and operations

## Why is it important?

- Industry standard for container orchestration
- Enables auto-scaling and self-healing
- Essential for production ML deployments

# The Scaling Problem

**Scenario:** Joe and all his friends request a prediction simultaneously

**Traditional approach:**

- Create a queue?
- Serve each request sequentially?
- Users wait... and wait... and wait...

## The Solution

**Scale up the number of computers!**

But who manages this scaling? **Kubernetes!**

## What Kubernetes provides:

- **Auto-scaling:** Automatically add/remove containers
- **Load balancing:** Distribute requests across instances
- **Self-healing:** Restart failed containers
- **Rolling updates:** Deploy new versions without downtime
- **Service discovery:** Apps find each other automatically

## Architecture components:

- Pods (smallest deployable units)
- Services (network access)
- Deployments (manage replicas)
- Ingress (reverse proxy/load balancer)

# GPU + Kubernetes: The Future

## Why GPUs + Kubernetes matter:

### Massive investments:

- \$500B on OpenAI Stargate project alone
- McKinsey: \$5.2 trillion needed for AI data centers by 2030
- OpenAI scaling to 7,500+ Kubernetes nodes

### Career implications:

- Growing demand for cloud infrastructure consultants
- Kubernetes + GPU expertise is highly valued
- Critical skill for modern ML engineering

### Source:

<https://openai.com/index/scaling-kubernetes-to-7500-nodes/>

## Essential GPU terminology for ML:

- **CUDA:** NVIDIA's parallel computing platform
- **Tensor Cores:** Specialized AI acceleration units
- **VRAM:** Video memory for model parameters
- **FP16/FP32:** Floating-point precision levels
- **Multi-GPU:** Training across multiple GPUs
- **GPU Memory:** Often the bottleneck in ML workloads

**Excellent resource:** Charles Frye's GPU Glossary

<https://modal.com/gpu-glossary>

*Credit: Charles Frye for this comprehensive resource*



**The complete journey:** PC  $\rightarrow$  Container  $\rightarrow$  Orchestration  $\rightarrow$  Cloud

**Key components we've covered:**

- **Docker:** Containerize your ML model
- **Kubernetes:** Orchestrate at scale
- **Cloud platforms:** AWS, GCP, Azure
- **Auto-scaling:** Handle variable loads

**Architecture pattern:**

- Load balancer/Reverse proxy
- Multiple container instances
- Auto-scaling based on demand
- Health checks and self-healing

# Deployment Architecture

## Typical production setup:

User Request → Load Balancer → Kubernetes Cluster → ML Model  
Pods

## Components:

- **Reverse Proxy:** Nginx, Traefik, or cloud load balancer
- **Ingress Controller:** Routes traffic to services
- **Service Mesh:** Advanced traffic management (optional)
- **Monitoring:** Prometheus, Grafana for observability

## Scaling triggers:

- CPU/Memory utilization
- Request queue length
- Custom metrics (inference time, etc.)

# Cloud Provider Options

## Major cloud platforms for ML deployment:

### AWS

- EKS (Kubernetes)
- SageMaker
- EC2 with GPUs
- Lambda (serverless)

### Google Cloud

- GKE (Kubernetes)
- Vertex AI
- Compute Engine
- Cloud Run

### Azure

- AKS (Kubernetes)
- Azure ML
- Virtual Machines
- Container Instances

## Key considerations:

- GPU availability and pricing
- Regional data requirements
- Integration with existing systems

## It's Coding Time!

Let's put theory into practice

### **What we'll implement:**

- Containerize an ML model with Docker
- Set up basic Kubernetes deployment
- Test scaling and load balancing
- Deploy to cloud (time permitting)

# Coding Session Structure

## Our hands-on approach:

- 1 Run the provided code examples
- 2 Read and understand what each part does
- 3 Discuss with your neighbor
- 4 Ask questions - to me, each other, or AI assistants
- 5 Experiment - modify and see what happens

## Resources Available

- Code examples in the repository
- Setup scripts for different operating systems
- ChatGPT/Claude for quick explanations
- Peer collaboration encouraged!

**Let's transform your ML model  
from a local script  
into a cloud-ready application!**

Questions before we start coding?

# What We've Covered

## Today's journey:

- **Introduction:** The ML-ops pipeline overview
- **Docker:** Containerization fundamentals
- **Kubernetes:** Orchestration and scaling
- **Cloud Deployment:** Production-ready architecture
- **Hands-on Coding:** From theory to practice

## Mission accomplished?

*“Acquire the knowledge to put a model from your PC into the cloud”*

**You now have the foundation!**

The building blocks are in place. Practice and iteration will make you proficient.

## Continue learning:

- **Docker:** Official documentation and tutorials
- **Kubernetes:** <https://kubernetes.io/docs/tutorials/>
- **GPU Glossary:** <https://modal.com/gpu-glossary> (Charles Frye)
- **Cloud platforms:** Provider-specific ML documentation
- **Practice:** Deploy your own models using today's concepts

## Next steps:

- Try the setup scripts for your OS
- Experiment with the provided code examples
- Join ML-ops communities and forums
- Consider cloud certifications (AWS, GCP, Azure)



Thank You!

# Thank You!

Questions & Discussion

**Contact:**

sveinung.myhre@example.com

**ReLU NTNU**

September 30, 2025

# References I